

Research Statment

Alex Sanchez-Stern

My research goal is to combine the techniques and challenges of programming languages research with those from adjacent fields like numerical analysis, security, and machine learning, to make it easier for programmers to write high-assurance software in a clear and concise manner. Research in the computer science literature has produced many theoretical techniques for producing high-assurance software. Unfortunately, many of these techniques are difficult to use, and require special training and intense manual labor to produce high-quality results. My research direction has been to combine existing approaches in new ways to bring the power of strong program reasoning to more programmers. In doing so, I've worked on tools to allow more programmers to write bug-free numerical code [14, 19](PLDI 2015, PLDI 2018), secure web-applications [?], and high-assurance verified software [17, 18], (CPP 2020, MAPL 2020).

1 Synthesizing and Debugging Numerical Software

Floating point math is a notoriously leaky abstraction over the real numbers, making writing reliable numerical code extremely difficult [8, 9, 20]. When the abstraction fails, rounding errors have caused the retraction of scientific articles [1, 3, 2], legal regulations in finance [6], and distorted stock market indices [13, 16]. Developers working with floating point often try to fix rounding error through ad-hoc methods, such as randomly perturbing the code until it works on test inputs [20, 9], or increasing the bit-width of their floating-point values, but these approaches can lead to further errors and significant performance degradation, and are completely ineffective for many problems. When development of numerical software has access to significant resources, developers can use formal numerical analysis to produce accurate programs, but these techniques require significant expertise, and the process is still slow and complicated.

I first addressed this challenge with Herbie, a tool I worked on with a team from the University of Washington. Herbie allows users to write down the real-number computation that they want to compute, and automatically transforms it into floating-point code which can compute it quickly and accurately. This allows users to program in abstractions they understand, and brings techniques of numerical analysis, once restricted to a few experts around the world, into the compilers of everyday programmers. Herbie is available as an official Racket package, a user-friendly web interface, and a plugin to the Glasgow Haskell Compiler. It has had patches accepted to widely used math libraries like MathJS, and has been used by researchers at a variety of institutions, including NASA.

However, in order to use Herbie, a user must know which segments of their code are non-trivial numerically, and worthy of analysis. This is not always an easy task, especially since much of the large numerical software in use is legacy software, written decades ago and maintained carefully. Extracting numerical code from such software for repair is extremely difficult, especially since the operations which interact to cause error can be distributed throughout the program source code, and can execute at very different points in the programs execution.

To bring the power of Herbie to such large software, I developed a binary analysis tool called Herbgrind. Built using the Valgrind framework, Herbgrind instruments binary programs to track the sources of floating-point error, and gather a variety of information about the error into an error report which can allow the user to easily use Herbie to improve the numerical code. Herbgrind uses a novel variety of anti-unification to synthesize program fragments representing the dynamic flow of values, allowing operations to be analyzed across data structures and control flow.

2 Languages for Database Security

The correctness of many web applications depends on enforcing rules about what information a user can access and provide during the course of the application. From password-leaking vulnerabilities, to SQL injection, to inferring the private personal information of another user, the consequences of improperly gated data access can be catastrophic. Unfortunately, the logic that enforces data access protection is often scattered throughout an application and duplicated, making it brittle to changing application requirements and features.

Promising work has been done developing policy languages, an approach where the developer writes down data access permissions in a centralized place, and it is enforced at every access point. This prevents bugs due to permission mismatches throughout the code. But it does not address the possibility that the policies written down are themselves buggy, nor does it address the needs of web applications where data schemas and policies are constantly changing.

Past policy languages have specified the policy of data at any given point of the development process, but not how they change as databases are migrated and policies are updated. These changes are important, as migration can leak secret data to the public, and introduce buggy policies.

To address this gap, I developed Watchdog with John Renner and a team at UCSD, a new kind of policy language. Watchdog allows the user to not just specify data layout and policies, but also data- and policy-migrations which transform the data and policies between versions. Additionally, our policy language is constructed around an SMT semantics which allows it to quickly check properties of large policies, like “is any data that was secret before the migration publicly accessible after?”.

3 Studying and Automating Proof Engineering

One of the most promising approaches to producing high-reliability software is *foundational verification*. In this approach, programmers write programs and verify properties of them in specialized tooling known as an “interactive theorem prover”, such as Coq [7] or Isabelle/HOL [15]. This approach has been used to verify a variety of software, including compilers [11], operating systems [10], database systems [12], file systems [5], distributed systems [21], and cryptographic primitives [4].

Unfortunately, the reliability of software produced in this manner comes at a high cost: expert labor. Even for those trained in the use of interactive theorem provers, proofs are extremely time-intensive to write. CompCert [11] took 6 person-years and 100,000 lines of Coq to write and verify, and seL4 [10], which is a verified version of a 10,000 line operating system, took 22 person-years to verify. And for those without a graduate degree in programming languages, these tools are generally inaccessible due to the grasp of concepts required.

So the question becomes, “how do we bring the benefits of foundational verification to a wider audience”? I started to answer this question with REPLICA, a collaboration with Talia Ringer at the University of Washington published at CPP 2020. With REPLICA we conducted the first user-study of proof engineers, in order to understand what kinds of changes they make to proofs during development, and where their time is spent. Using the data from this study, we could understand the needs of proof engineers, in order to make the field more widely accessible to programmers through tooling.

From the data gathered through REPLICA, I moved on to tackling one of the most pressing problems in the accessibility of proof engineering: the amount of labor required to write and maintain large numbers of proofs, many of which are re-expressing simple facts about custom data structures. With Proverbot9001, published at MAPL2020, my team built a system to tackle this problem; a tool which uses machine learning to attempt to build proofs of arbitrary logical statements. Unlike other work which takes a low-level approach to building the syntactic constructs of the proofs, Proverbot9001 makes use of the tactic systems present in modern proof assistants, where small self-contained search procedures are invoked in sequence to construct a final proof term. Proverbot9001 learns from existing proofs to determine likely-useful tactics, and uses this knowledge to guide a search tree which constructs the final proof. With this approach, Proverbot9001 is able to re-discover almost a quarter of proofs in the CompCert verified C compiler, drastically reducing the amount of human labor required to maintain such a projects.

4 Future Work

My work so far has only begun to scratch the surface of what’s possible, both in bringing Programming Language tools to new domains, and in using cutting edge technologies like deep neural networks to solve the centuries-old problem of proof synthesis. As the amount of software in our world grows exponentially, so too does the need for reliability in that software; however we are still broadly under-equipped to deliver the reliability which our new world demands. With software permeating ever-more personal aspects of our

lives, the safeguarding of our digital data from malicious actors relies on being able to trust the software which we run.

The goal of my future work broadly is to expand the feasibility of proof synthesis in the software development practice. This includes developing new techniques and models for supervised learning of proofs, as well as expanding proof-synthesis systems to include reinforcement learning and other innovations from the machine learning literature. It also includes developing new tooling for bringing this proof synthesis to developers, including prediction-guided interfaces for developers writing more complex proofs, and new language tooling to bring the reliability of verified software into the languages that developers use most often.

The application of new machine-learning techniques to the problem of proof-synthesis is an exciting new direction, only beginning to be explored by the Programming Languages community. A multitude of approaches have emerged, but they are just beginning to scratch the surface of what is possible. In recent years the literature on graph neural networks has begun to blossom, and with it new opportunities for encoding program and proof syntax into feature spaces for deep neural networks to process. Additionally, the development of AlphaZero at Google’s DeepMind laboratory has shown that reinforcement learning in structured environments can outpace even supervised learning from large amounts of expert data; my future work will explore applying this insight to proof-synthesis, another highly structured environment.

But even with improving proof-synthesis technology, there is still a large gap between the specialized proof assistants in which verified software is produced, and the tooling in which most software is developed. Attempts to bridge this gap, like the Dafny programming language, are promising but still outside of the reach of most developers. In my future work, I intend to explore the connection between the invariants and proof approaches required to synthesize a proof of program correctness, and the developer intuition for correctness, which can and often is written down in informal comments accompanying the code. By bringing together these two tasks, I hope to decrease the developer labor required to verify software while at the same time increasing the detail and reliability of code documentation.

I also plan to continue my work collaborating with my colleagues in other fields to bring programming language techniques into the problems faced by developers in a variety of domains. My experience with Herbie, Herbgrind, and Watchdog has shown me the wealth of opportunity in collaborations across domains to produce high-impact tooling. While these collaborations will depend highly on the colleagues and expertise I have available, I’m particularly interested in continuing my work applying Programming Language techniques to security domains, and exploring the interaction between tooling and machine learning from new perspectives.

References

- [1] M. Altman, J. Gill, and M. P. McDonald. *Numerical Issues in Statistical Computing for the Social Scientist*. Springer-Verlag, 2003.
- [2] M. Altman and M. McDonald. The robustness of statistical abstractions. *Political Methodology*, 1999.
- [3] M. Altman and M. P. McDonald. Replication with attention to numerical accuracy. *Political Analysis*, 11(3):302–307, 2003.
- [4] A. W. Appel. Verification of a cryptographic primitive: Sha-256. *ACM Trans. Program. Lang. Syst.*, 37(2):7:1–7:31, Apr. 2015.
- [5] H. Chen, T. Chajed, A. Konradi, S. Wang, A. İleri, A. Chlipala, M. F. Kaashoek, and N. Zeldovich. Verifying a high-performance crash-safe file system using a tree specification. In *Proceedings of the 26th Symposium on Operating Systems Principles, SOSP ’17*, pages 270–286, New York, NY, USA, 2017. ACM.
- [6] European Commission. *The introduction of the euro and the rounding of currency amounts*. Euro papers. European Commission, Directorate General II Economic and Financial Affairs, 1998.

- [7] J.-C. Filliâtre, H. Herbelin, B. Barras, B. Barras, S. Boutin, E. Giménez, S. Boutin, G. Huet, C. Muñoz, C. Cornes, C. Cornes, J. Courant, J. Courant, C. Murthy, C. Murthy, C. Parent, C. Parent, C. Paulin-mohring, C. Paulin-mohring, A. Saibi, A. Saibi, B. Werner, and B. Werner. The coq proof assistant - reference manual version 6.1. Technical report, 1997.
- [8] W. Kahan. Miscalculating area and angles of a needle-like triangle. Technical report, University of California, Berkeley, Mar. 2000.
- [9] W. Kahan and J. D. Darcy. How Java’s floating-point hurts everyone everywhere. Technical report, University of California, Berkeley, June 1998.
- [10] G. Klein, K. Elphinstone, G. Heiser, J. Andronick, D. Cock, P. Derrin, D. Elkaduwe, K. Engelhardt, R. Kolanski, M. Norrish, T. Sewell, H. Tuch, and S. Winwood. sel4: Formal verification of an os kernel. In *Proceedings of the ACM SIGOPS 22Nd Symposium on Operating Systems Principles*, SOSP ’09, pages 207–220, New York, NY, USA, 2009. ACM.
- [11] X. Leroy. Formal verification of a realistic compiler. *Communications of the ACM*, 52(7):107–115, 2009.
- [12] G. Malecha, G. Morrisett, A. Shinnar, and R. Wisnesky. Toward a verified relational database management system. In *Proceedings of the 37th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*, POPL ’10, pages 237–248, New York, NY, USA, 2010. ACM.
- [13] B. D. McCullough and H. D. Vinod. The numerical reliability of econometric software. *Journal of Economic Literature*, 37(2):633–665, 1999.
- [14] P. Panckhka, A. Sanchez-Stern, J. R. Wilcox, and Z. Tatlock. Automatically improving accuracy for floating point expressions. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI ’15. ACM, 2015.
- [15] L. C. Paulson. Natural deduction as higher-order resolution. *CoRR*, cs.LO/9301104, 1993.
- [16] K. Quinn. Ever had problems rounding off figures? This stock exchange has. *The Wall Street Journal*, page 37, November 8, 1983.
- [17] T. Ringer, A. Sanchez-Stern, D. Grossman, and S. Lerner. Replica: Repl instrumentation for coq analysis. In *Proceedings of the 9th ACM SIGPLAN International Conference on Certified Programs and Proofs*, CPP 2020, page 99–113, New York, NY, USA, 2020. Association for Computing Machinery.
- [18] A. Sanchez-Stern, Y. Alhessi, L. Saul, and S. Lerner. Generating correctness proofs with neural networks. In *Proceedings of the 4th ACM SIGPLAN International Workshop on Machine Learning and Programming Languages*, MAPL 2020, page 1–10, New York, NY, USA, 2020. Association for Computing Machinery.
- [19] A. Sanchez-Stern, P. Panckhka, S. Lerner, and Z. Tatlock. Finding root causes of floating point error. *SIGPLAN Not.*, 53(4):256–269, June 2018.
- [20] N. Toronto and J. McCarthy. Practically accurate floating-point math. *Computing in Science Engineering*, 16(4):80–95, July 2014.
- [21] J. R. Wilcox, D. Woos, P. Panckhka, Z. Tatlock, X. Wang, M. D. Ernst, and T. Anderson. Verdi: A framework for implementing and formally verifying distributed systems. In *Proceedings of the 36th ACM SIGPLAN Conference on Programming Language Design and Implementation*, PLDI ’15, pages 357–368, New York, NY, USA, 2015. ACM.