



**Alex Sanchez-Stern**, Abhishek Varghese, Zhanna Kaufman,  
Dylan Zhang, Talia Ringer, Yuriy Brun



Automating Formal Verification with  
Reward-Free Reinforcement Learning



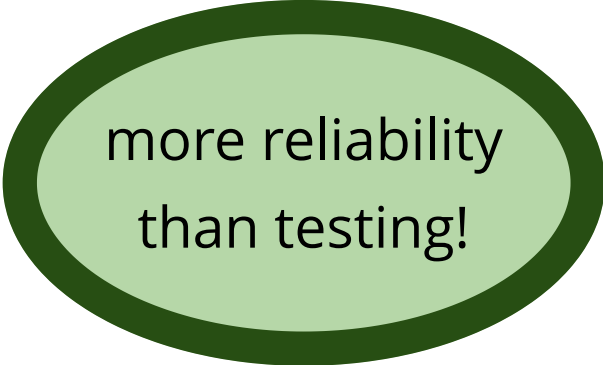
Automating Formal Verification with  
Reward-Free Reinforcement Learning

# Formal Verification

writing proofs about programs!

# Formal Verification

writing proofs about programs!



more reliability  
than testing!



# Formal Verification

writing proofs about programs!

more reliability  
than testing!

labor  
intensive



# Formal Verification

writing proofs about programs!

more reliability  
than testing!




labor  
intensive

requires specialized  
expertise




# Automating Formal Verification





# Automating Formal Verification


Computational  
Logic



# Automating Formal Verification

Computational  
Logic

Tree  
Search





# Automating Formal Verification

Computational  
Logic


Machine  
Learning

Tree  
Search








Proof Contexts



Actions







Unbounded Search Space

There are some techniques that can  
help us prune the search tree







There are some techniques that can  
help us prune the search tree




It's hard to explore when you don't know  
where you are!




We can search more  
efficiently if we can evaluate  
proof states




We can search more  
efficiently if we can evaluate  
proof states




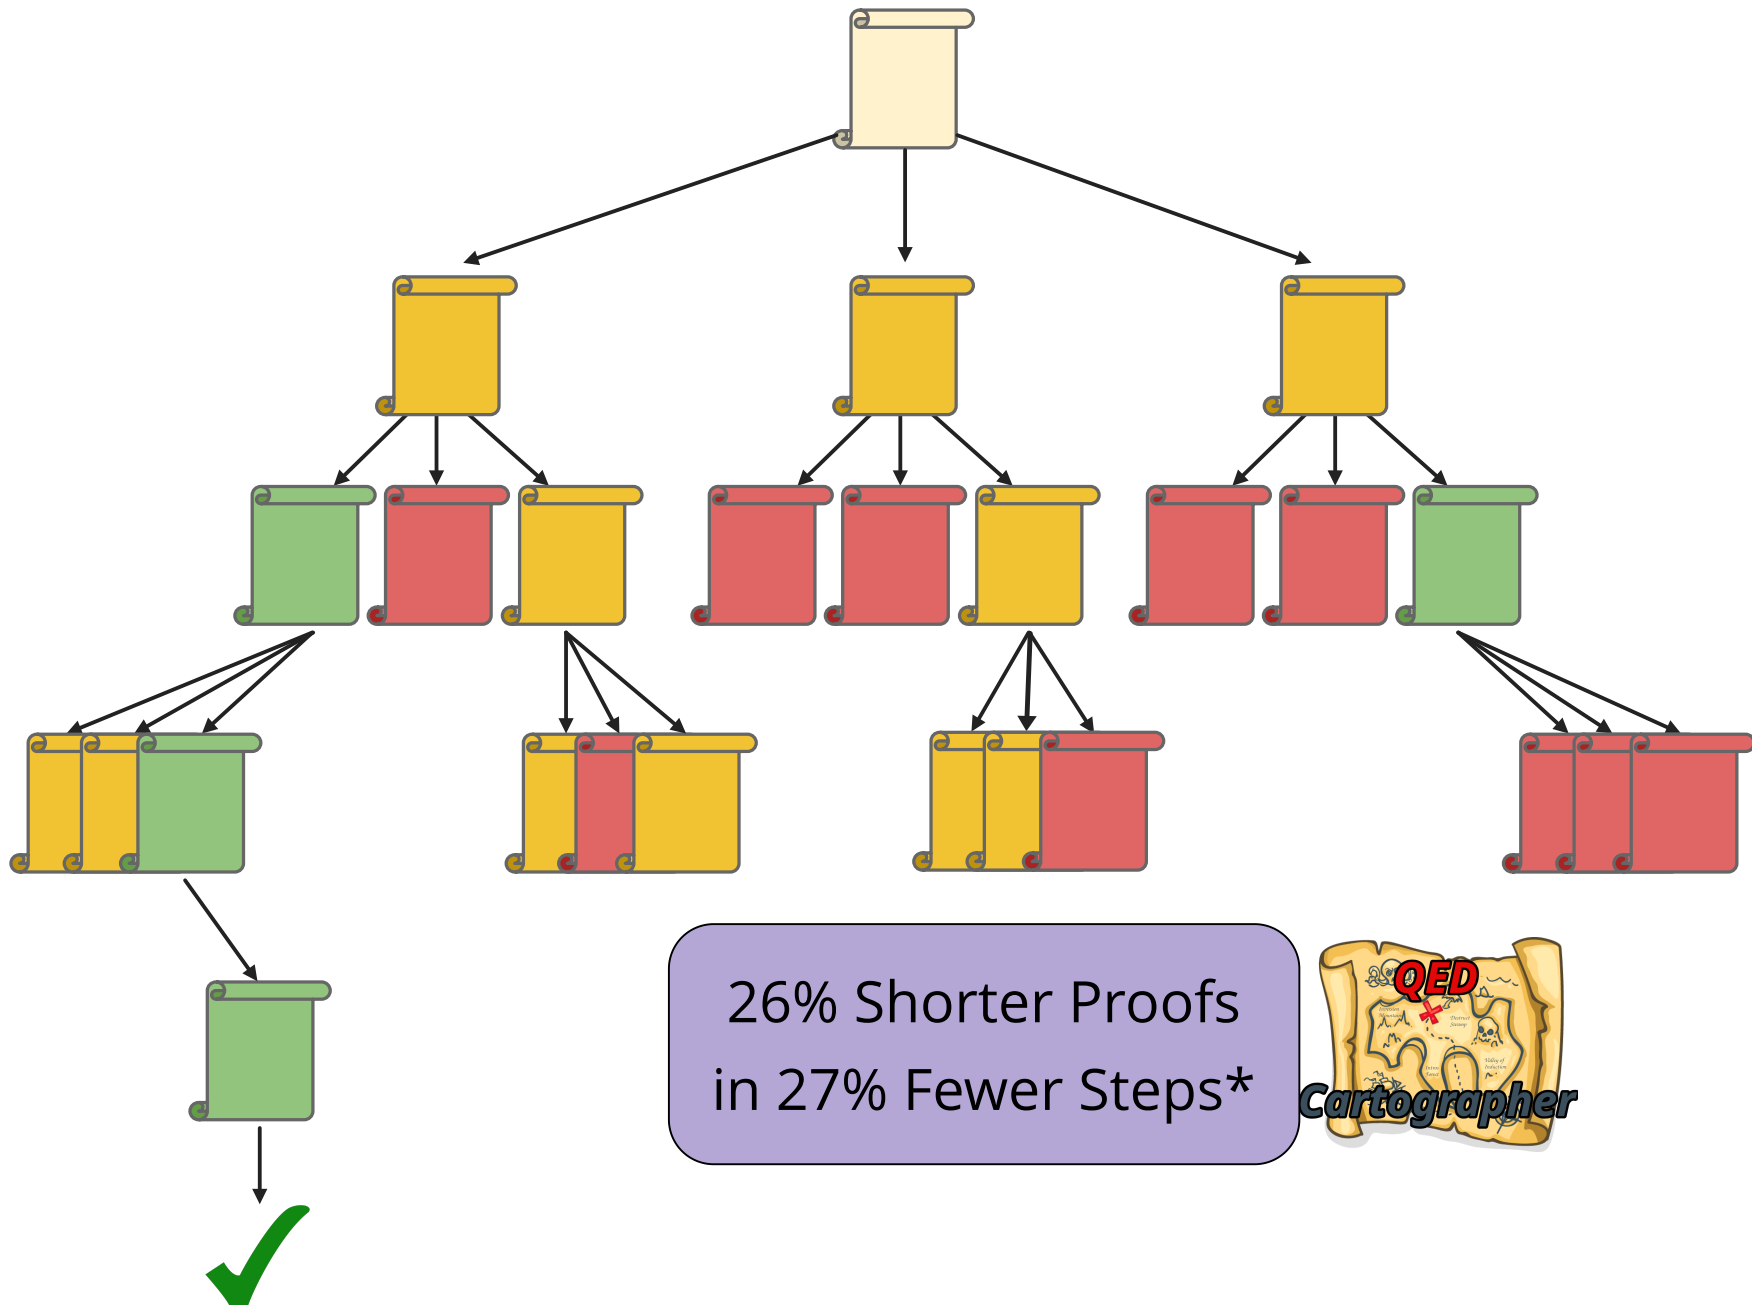
We can search more  
efficiently if we can evaluate  
proof states

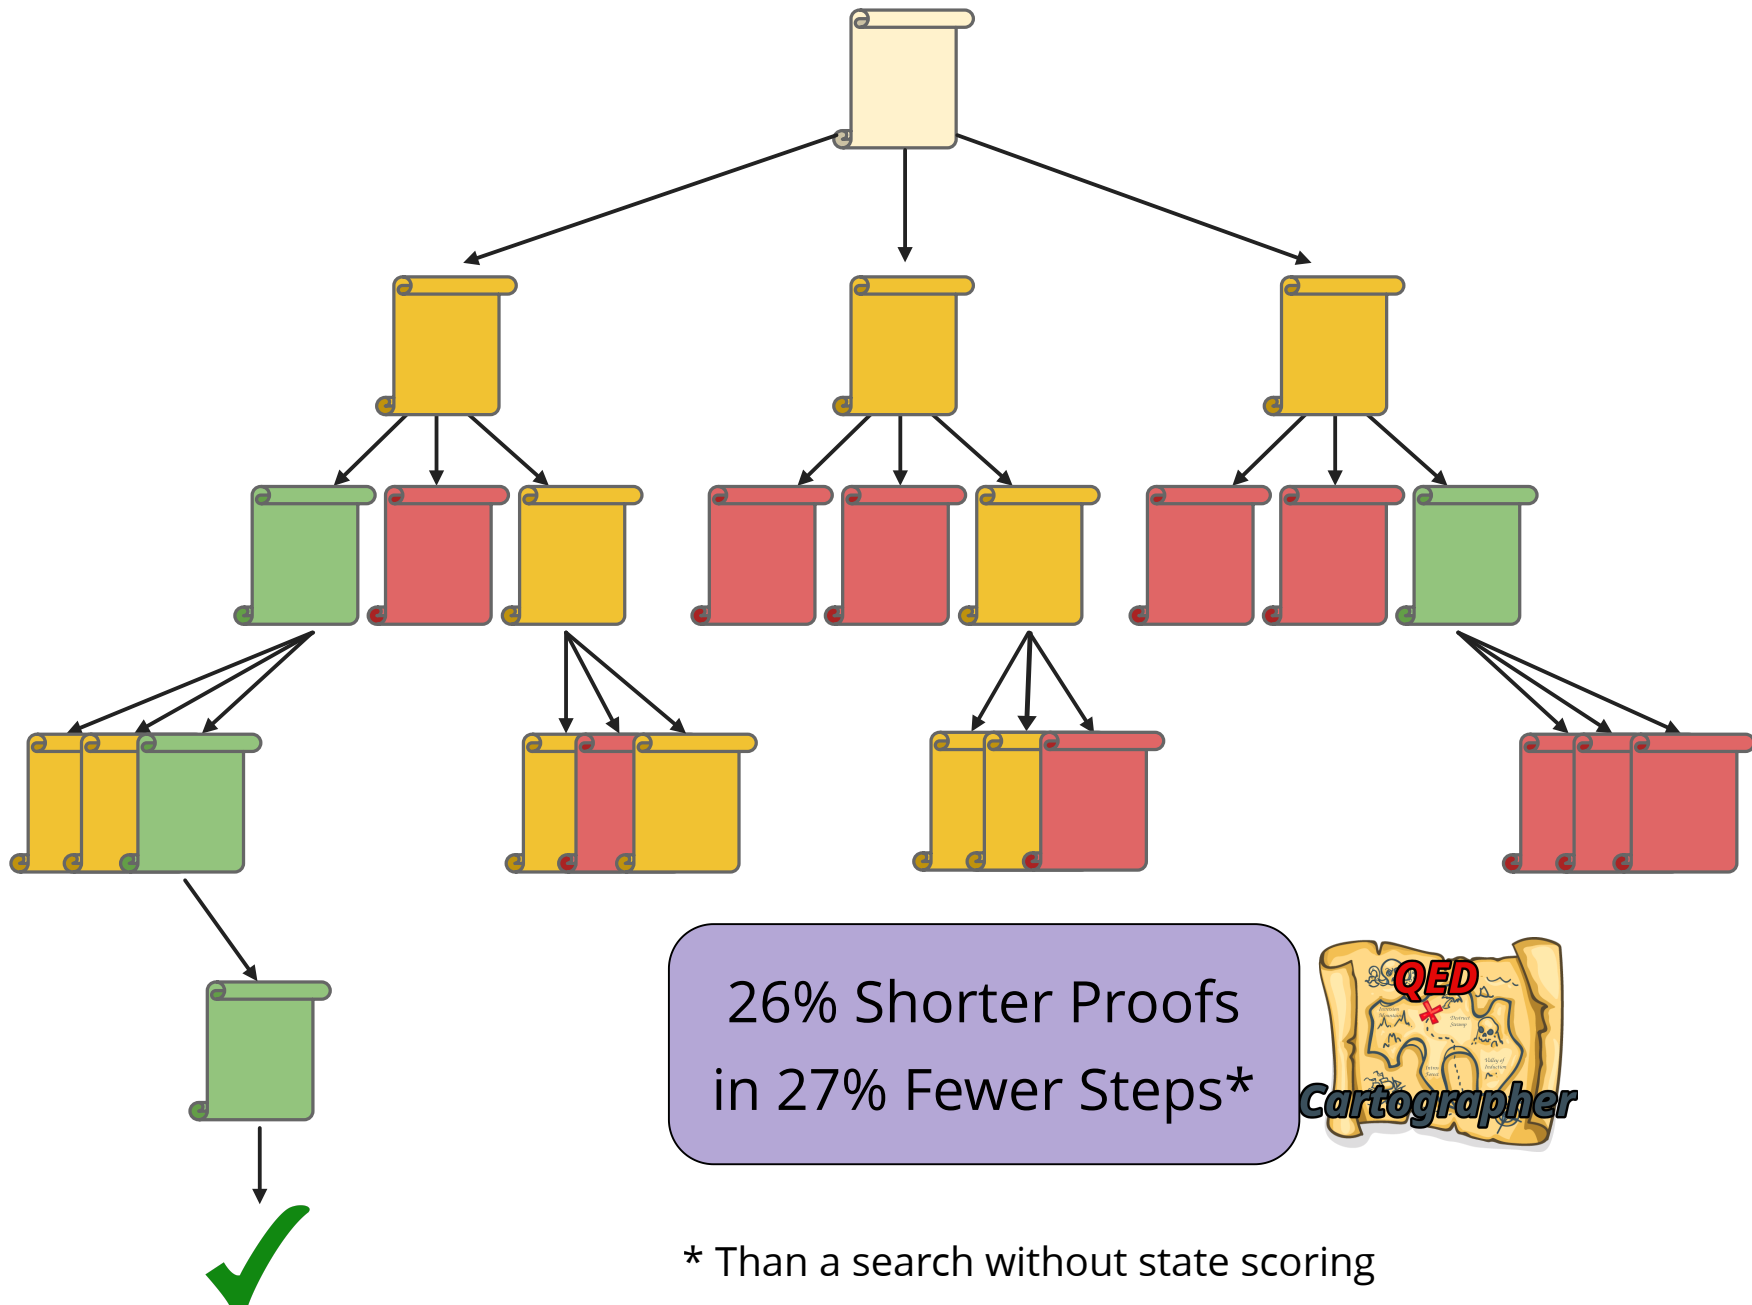


We can search more  
efficiently if we can evaluate  
proof states














Predictor




Search




Theorem Prover





Predictor



Search



Theorem Prover



"Reward-free Reinforcement Learning"

"Reward-free Reinforcement Learning"




In particular, V-learning

# This Talk


# This Talk

V-Learning




# This Talk

V-Learning




Limitations in Proofs




# This Talk


V-Learning



Limitations in Proofs




Adapting to Proofs







# Classic V-Learning



# Classic V-Learning




# Classic V-Learning



$$V(S1) = \max(\begin{aligned} &\text{penny} + V(S2), \\ &(\text{quarter} + V(S4)), \\ &\text{nickel} + V(S5)) \end{aligned}$$

# Classic V-Learning



$$V(S1) = \max(\text{penny} + \gamma V(S2), \gamma(\text{quarter} + \gamma V(S4)), \text{dime} + \gamma V(S5))$$

$\gamma$  = time discount

finding rewards sooner is better!

$$V(S1) = \max(\text{penny} + \gamma V(S2), \gamma (\text{quarter} + \gamma V(S4)), \text{nickel} + \gamma V(S5))$$


$$V(S1) = \max(\text{👤} + \gamma V(S2), \gamma (\text{👤}) + \gamma V(S4), \text{👤} + \gamma V(S5))$$

Generalizes to




$$V(S) = \max_{a \in \text{actions}(S)} (R(S, a) + \gamma V(\text{next-state}(S, a)))$$

# V-Learning in Practice: Iterative Updates



# V-Learning in Proofs:


## The Sparse Reward Problem






# V-Learning in Proofs:

## The Sparse Reward Problem



# V-Learning in Proofs:

## The Sparse Reward Problem






**Insight:**

**Proofs have Useful Structure**

# Abstraction



**State**



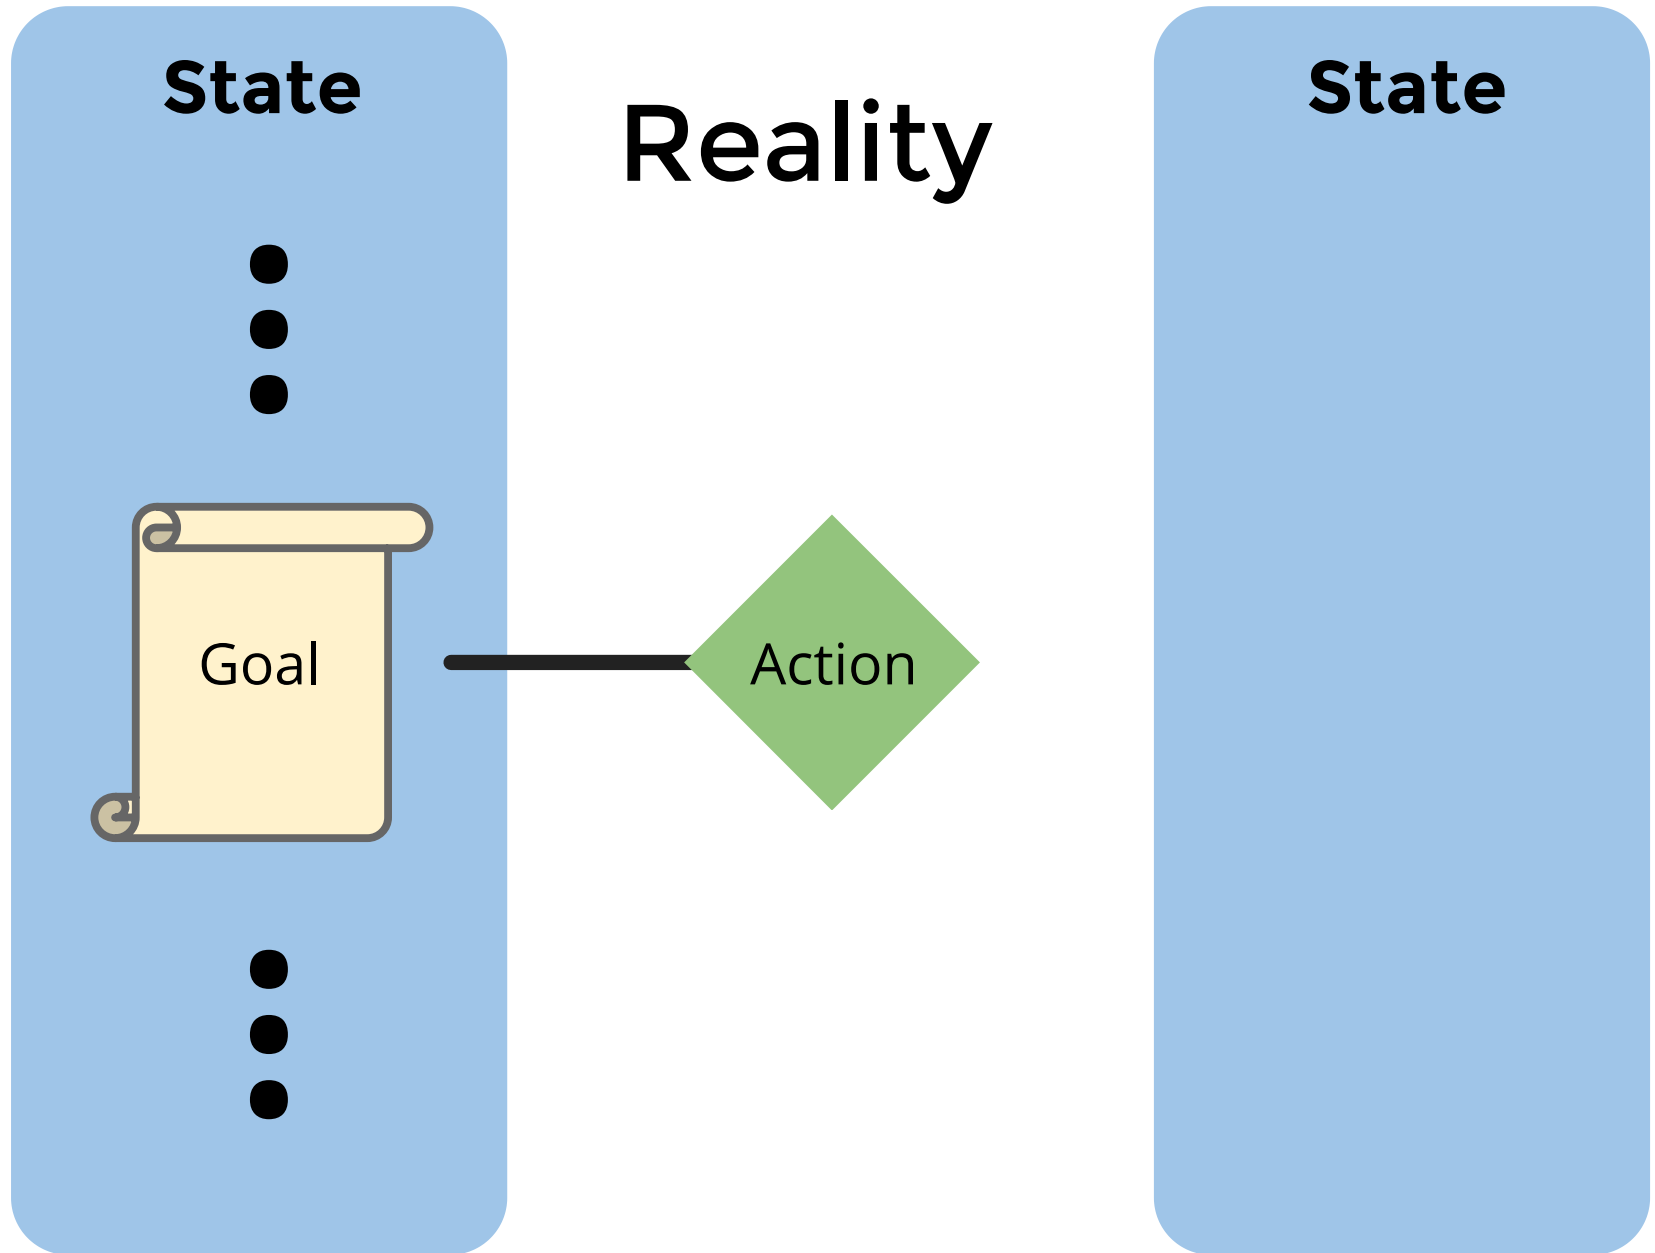
Goal

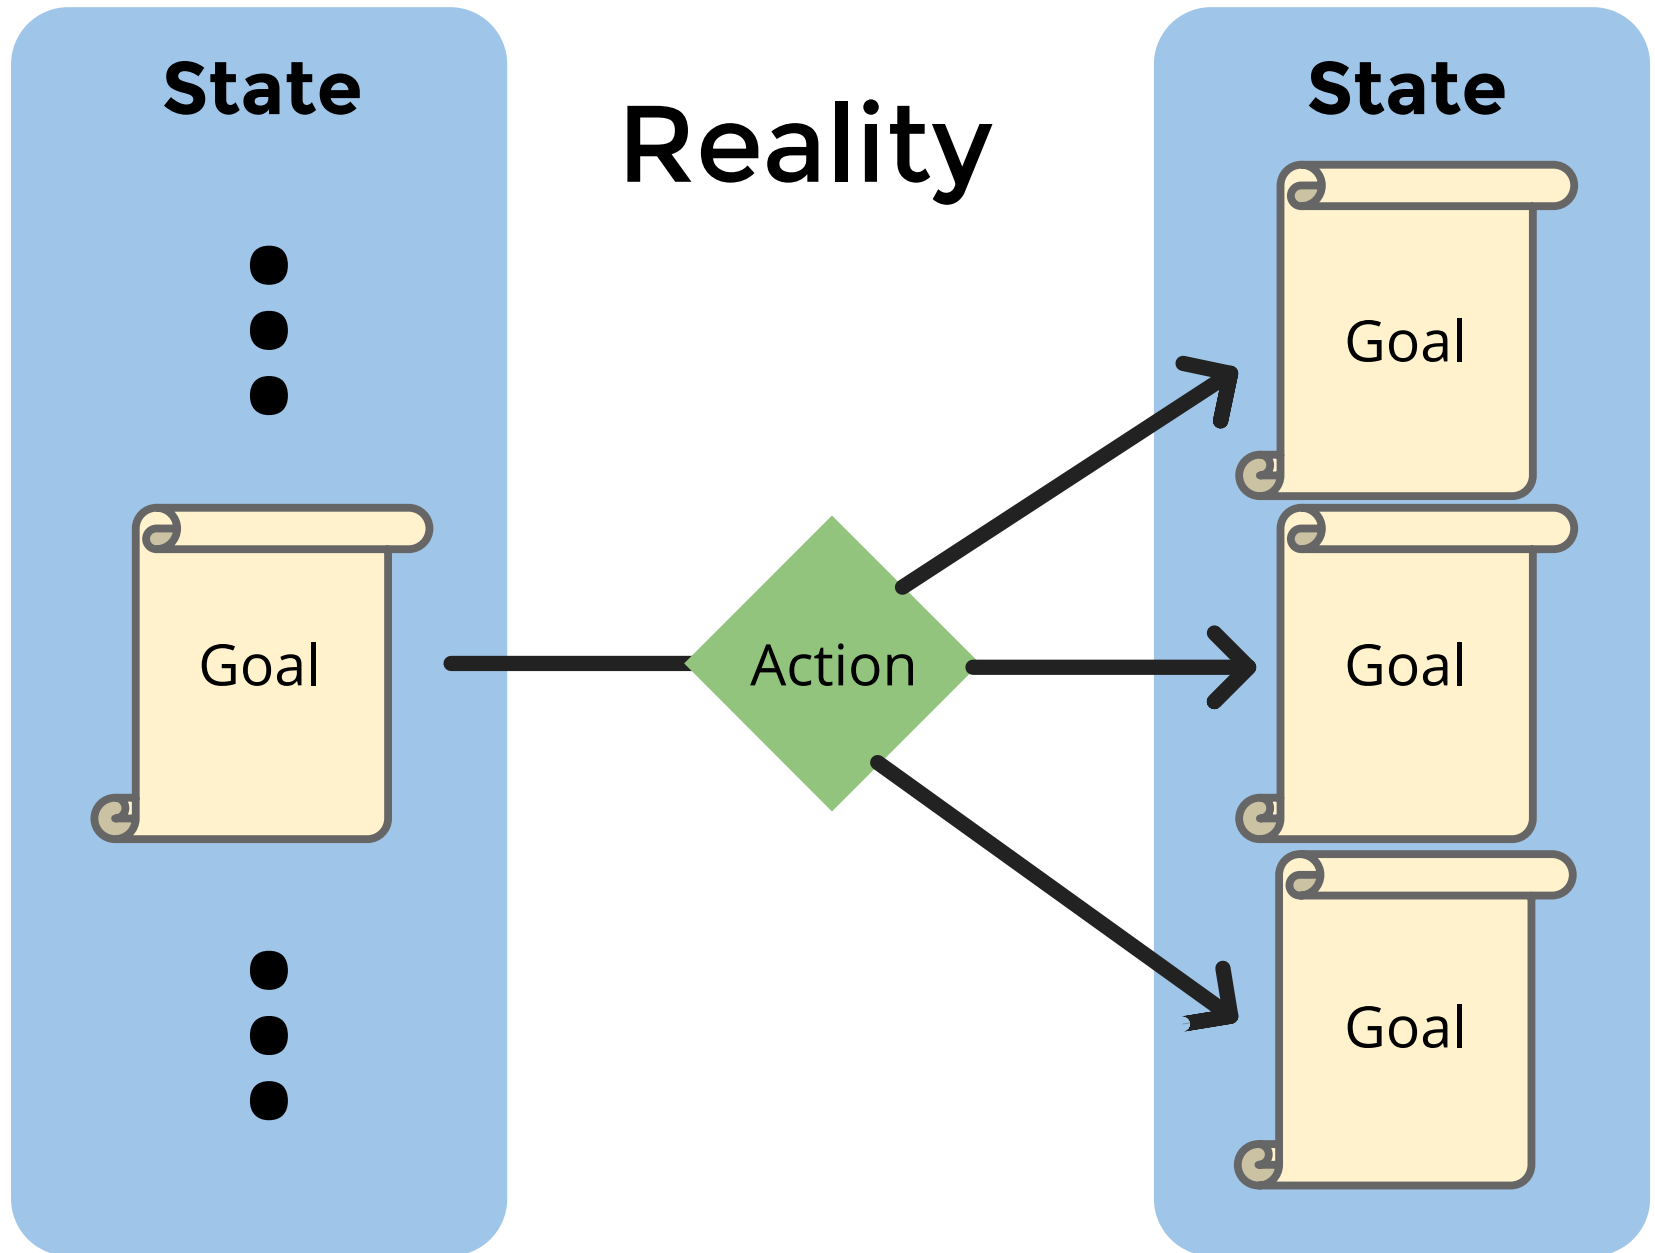


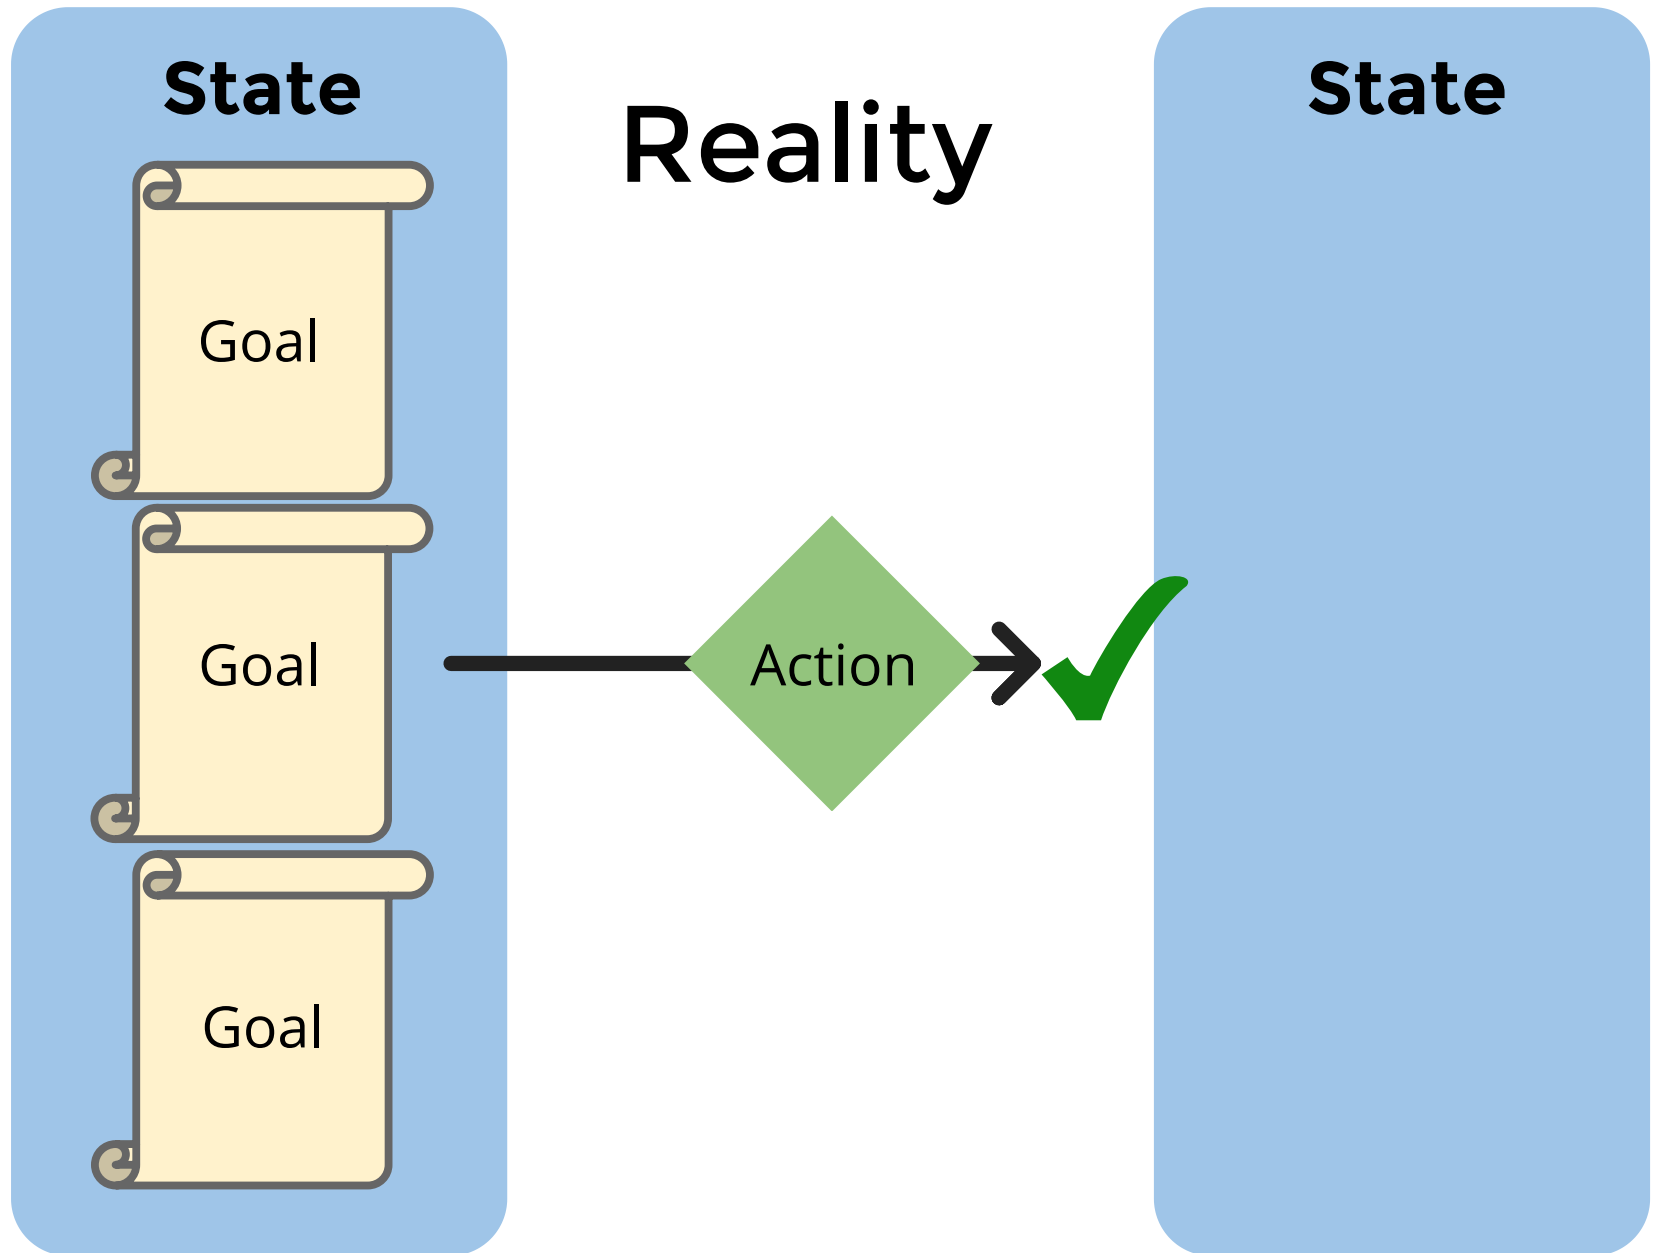
**Reality**

Action

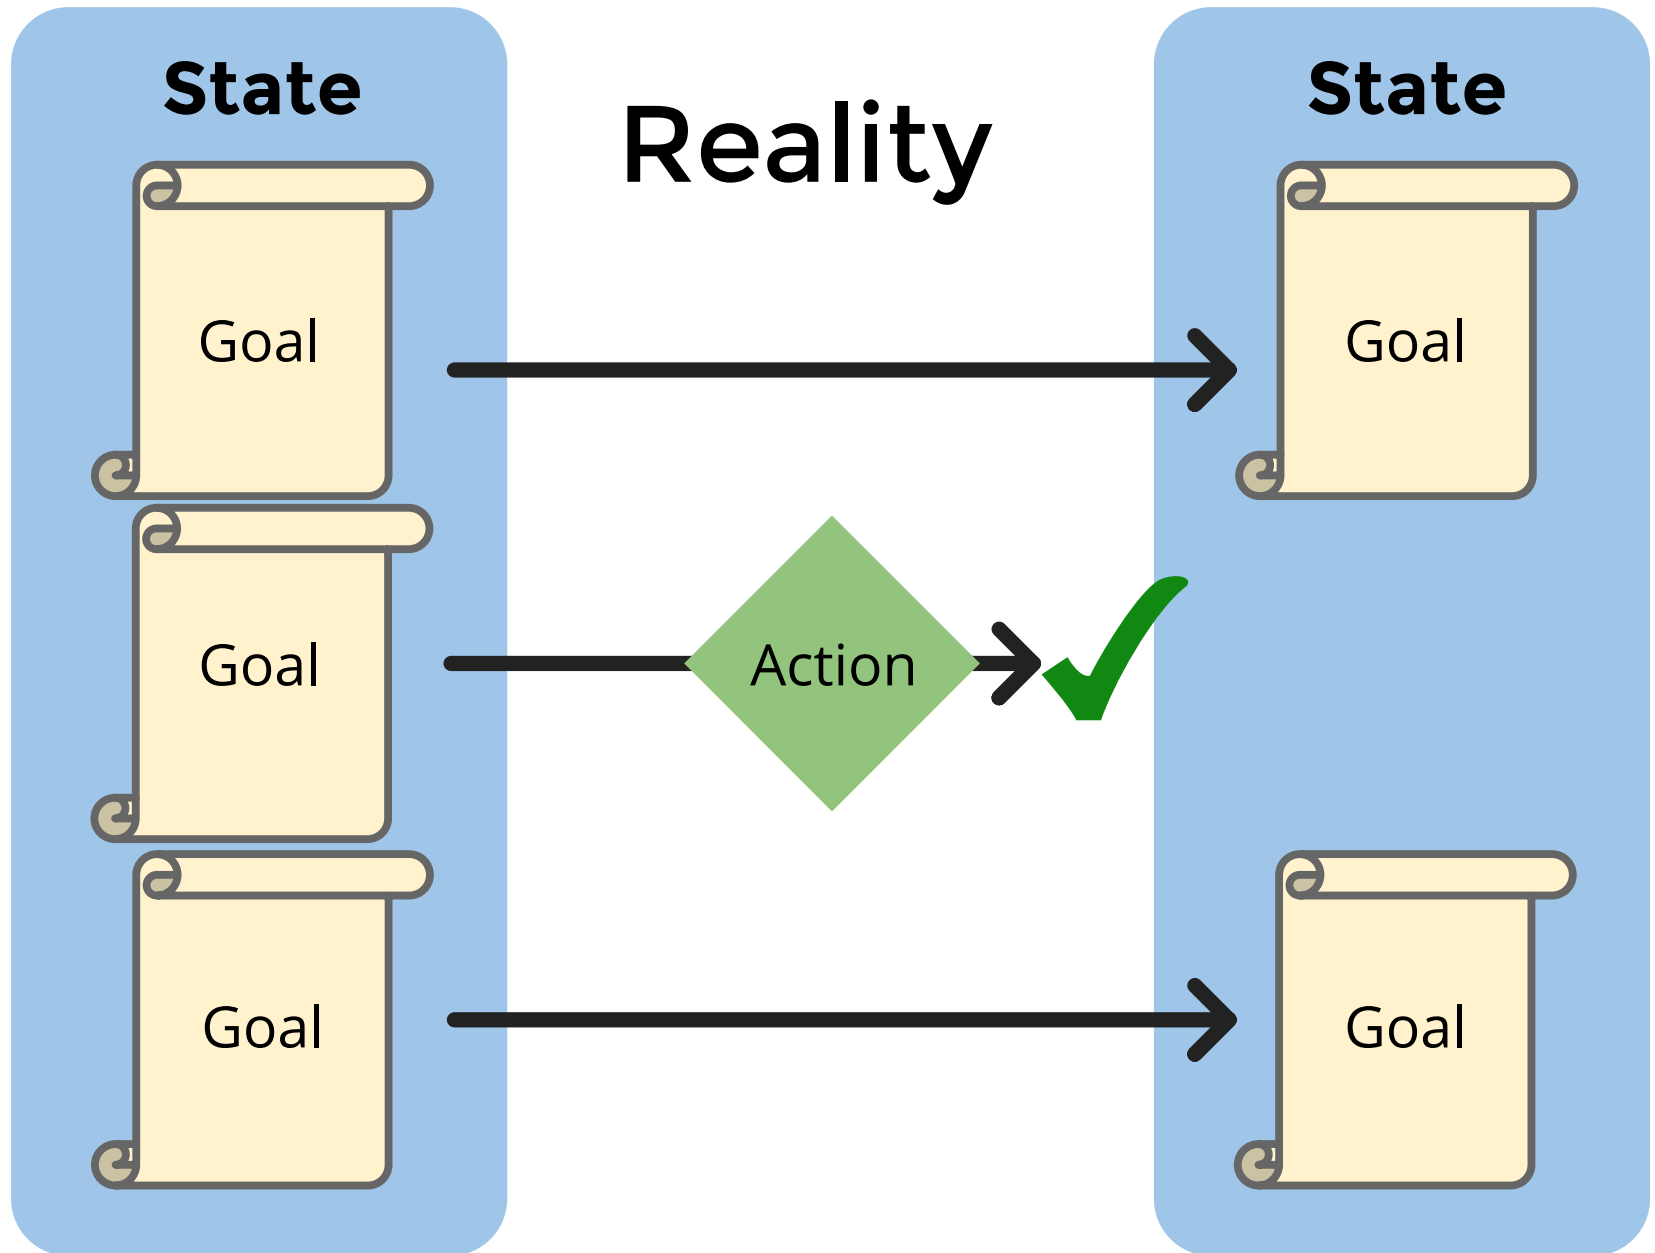
**State**













Adding extra rewards can lead to bad behavior

Adding extra rewards can lead to bad behavior




Adding extra rewards can lead to bad behavior



# Adding extra rewards can lead to bad behavior


Inductive Case



Base Case

# Adding extra rewards can lead to bad behavior

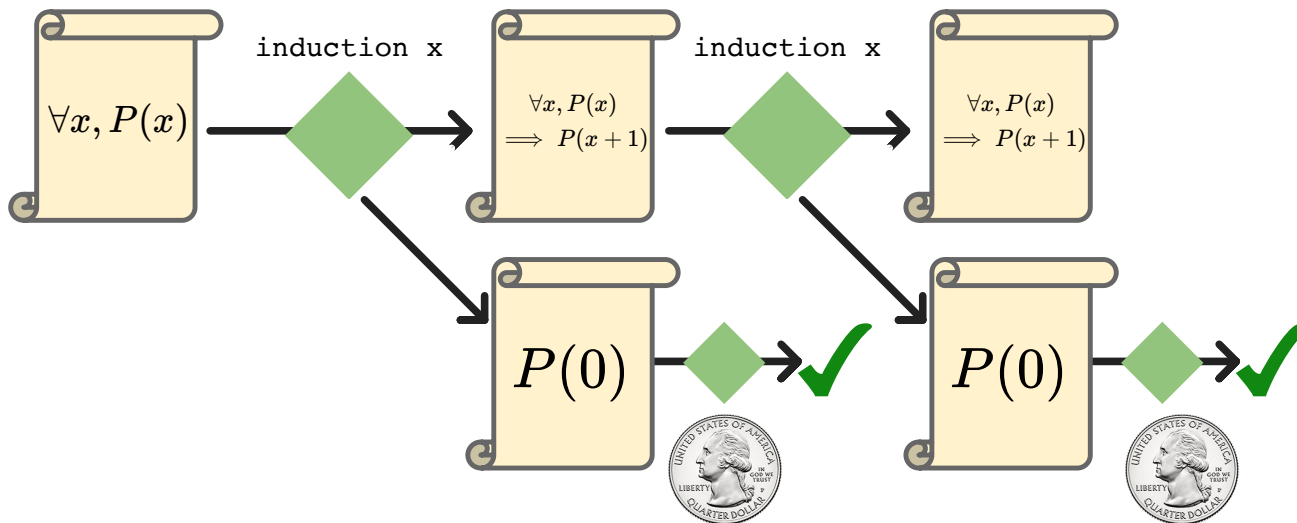
Inductive Case



Base Case

# Adding extra rewards can lead to bad behavior

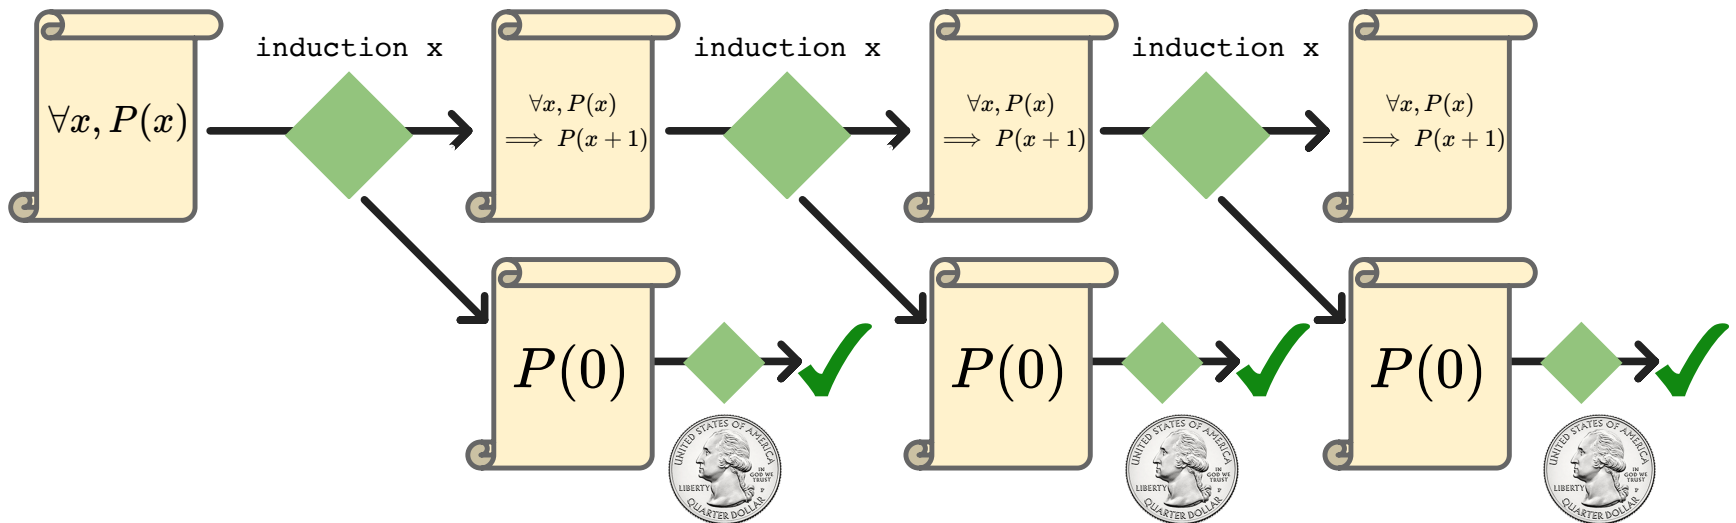
Inductive Case



Base Case

# Adding extra rewards can lead to bad behavior

Inductive Case

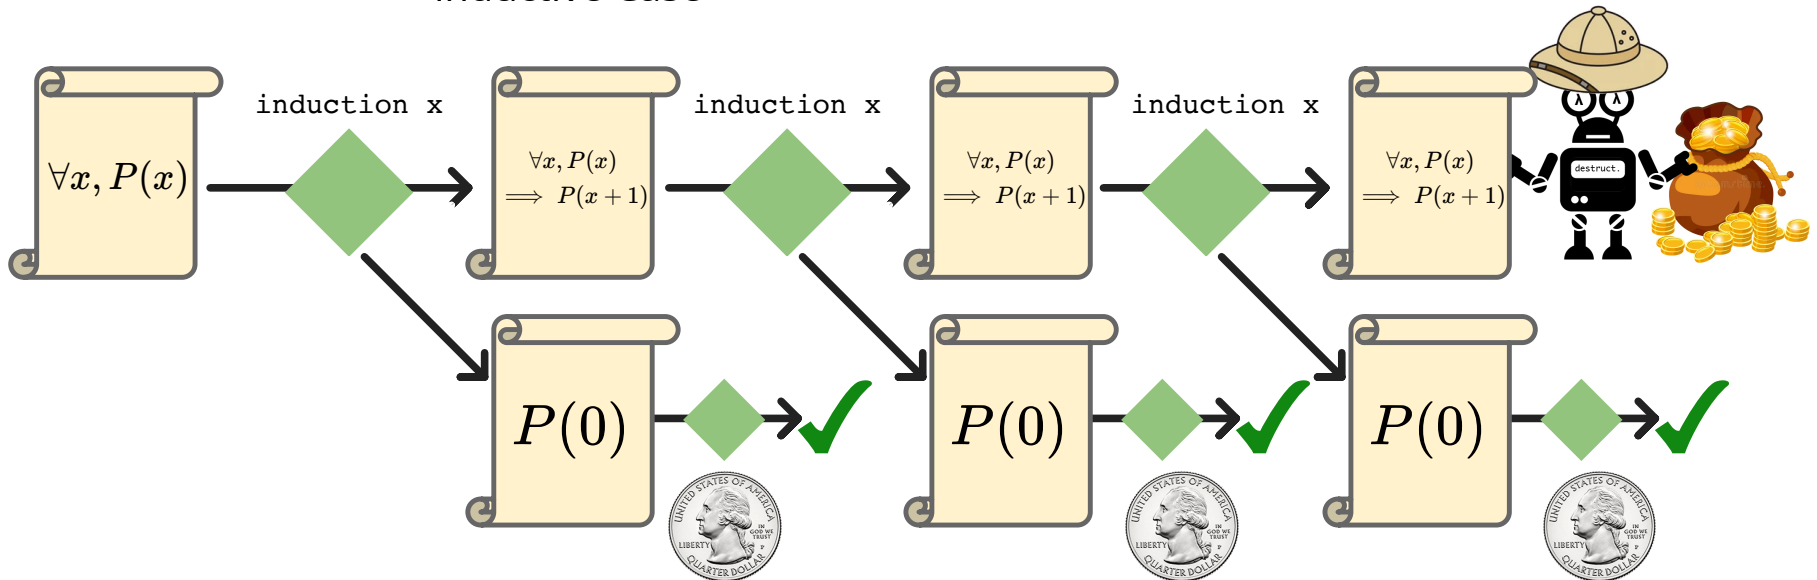


Base Case



# Adding extra rewards can lead to bad behavior


Inductive Case



Base Case

# Adding extra rewards can lead to bad behavior

Inductive Case



Base Case

Reward-free doesn't have this problem!


# What We Need

A new update equation that accounts for  
the branching structure of proofs

## Assumptions:


- The state of a completed proof has value 1

$$V(\checkmark) = 1$$



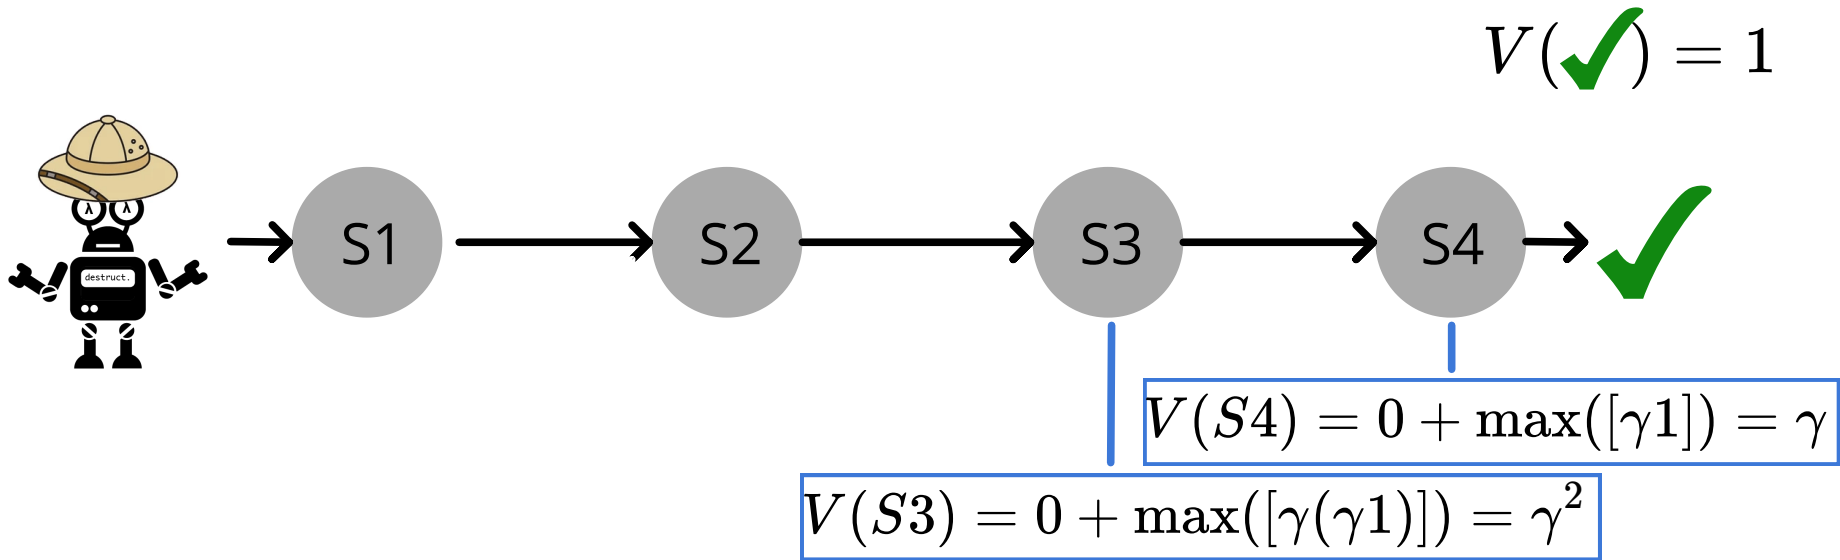
## Assumptions:

- The state of a completed proof has value 1



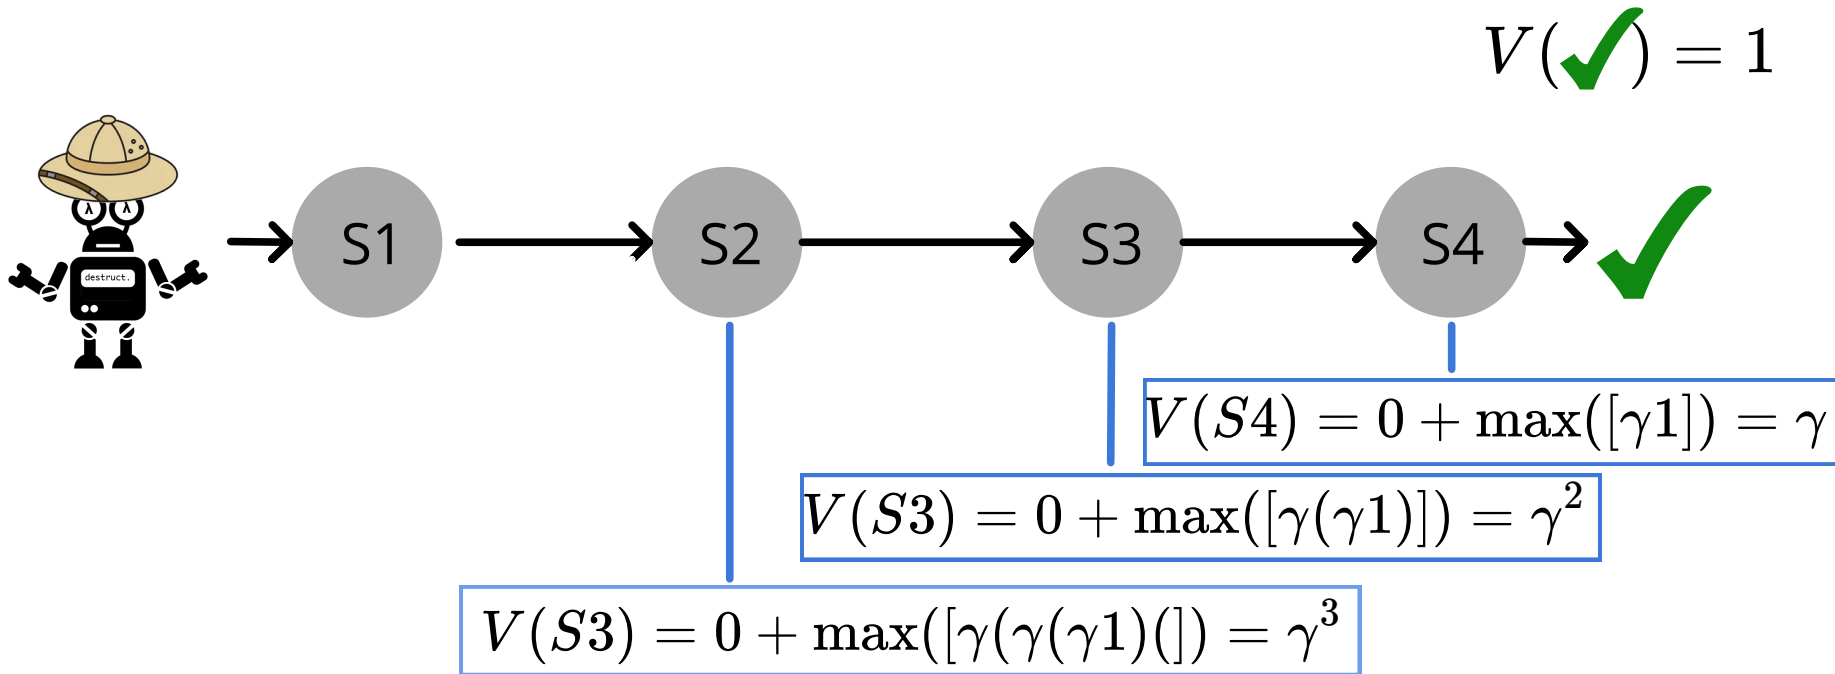
## Assumptions:

- The state of a completed proof has value 1



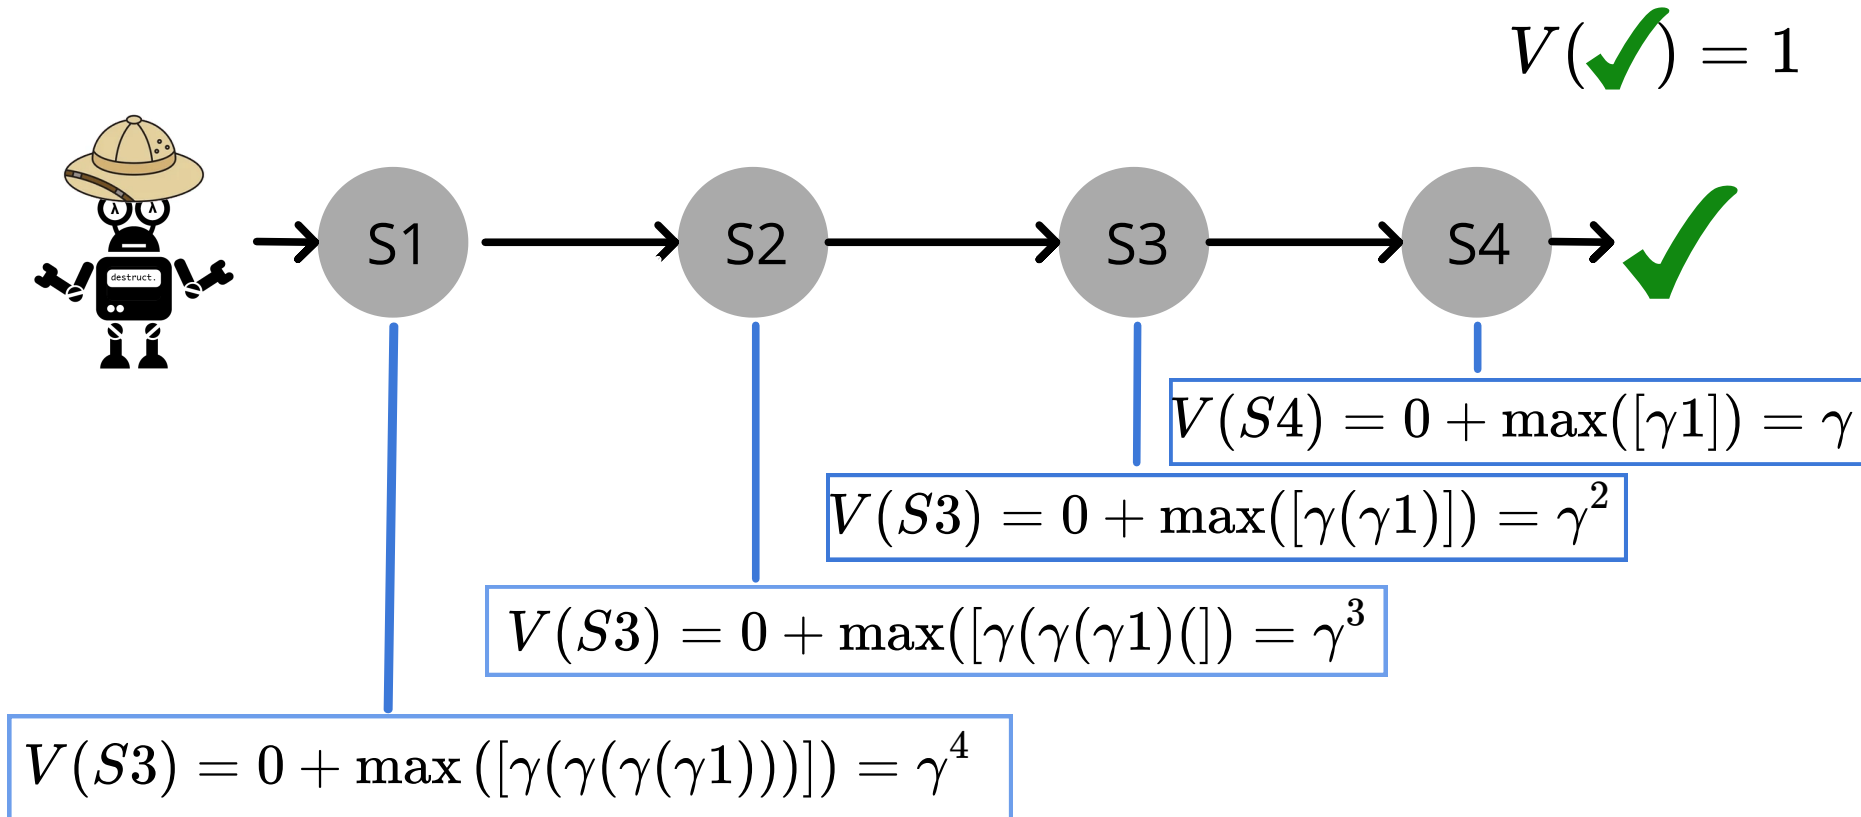
## Assumptions:

- The state of a completed proof has value 1





## Assumptions:

- The state of a completed proof has value 1






Assumption: The state of a completed proof  
has value 1





$$V(\checkmark) = 1$$

$$V(S) = \gamma^{(\text{number of steps left})}$$




$V(G1) = ???$






$$V(G1) = ???$$



$m$  = Steps to complete proof from G2

$n$  = Steps to complete proof from G3



$$V(G1) = ???$$

G1


$m$  = Steps to complete proof from G2

$n$  = Steps to complete proof from G3

Steps to complete proof from G1 =  $m + n + 1$

G2

G3



$$V(G1) = ???$$

G1

$m$  = Steps to complete proof from G2

$n$  = Steps to complete proof from G3

Steps to complete proof from G1 =  $m + n + 1$

$$V(G1) = \gamma * V(G2) * V(G3)$$

G2

G3

## Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$

## Update Equation for Branch-Structured Proofs

$$V(G) = \boxed{\max_{a \in \text{actions}(G)}} \left( \gamma \prod_{G' \in \text{next-states}(G, a)} V(G') \right)$$

## Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$



## Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \boxed{\prod_{G' \in \text{next-states}(G,a)} V(G')} \right)$$

## Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$

# Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$


## Update Equation for Branch-Structured Proofs

$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$



## Update Equation for Branch-Structured Proofs


$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$




# Don't need one!



# Don't need one!




# Don't need one!



$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$




# Don't need one!




$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$
$$\gamma \prod_{G' \in \emptyset} V(G')$$


# Don't need one!



$$V(G) = \max_{a \in \text{actions}(G)} \left( \gamma \prod_{G' \in \text{next-states}(G,a)} V(G') \right)$$
$$\gamma \prod_{G' \in \emptyset} V(G')$$
$$\gamma(1)$$




Automating Formal Verification with  
Reward-Free Reinforcement Learning



Automating Formal Verification with  
Reward-Free Reinforcement Learning

26% Shorter Proofs  
in 27% Fewer Steps

# Benchmark: CoqGym




124 Coq Projects

68,501 Theorems

85/15 train-test split

# Baseline: Proverbot9001 (updated)


QEDCartographer, except without state  
scoring-based search




# Baseline: Proverbot9001 (updated)


QEDCartographer, except without state  
scoring-based search


Uses a variant of depth-first search instead



















26% Shorter Proofs  
than Proverbot9001



Shorter proofs are:

- More readable
- More maintainable
- Less expensive to check





Proves slightly more theorems, and proves complementary theorems



## Automating Formal Verification with Reward-Free Reinforcement Learning

Uses a new V-value equation for branching goal structure

Makes producing verified-correct code easier and faster

Preprint available at [alexsanchezstern.com](http://alexsanchezstern.com)